

APS-DFD meeting, November 2007

# **Analysis of Turbulence Datasets using a Database Cluster: Requirements, Design, and Sample Applications**

**ITR group @ JHU**

Charles Meneveau<sup>1</sup>, Yi Li<sup>1\*</sup>, Eric Perlman<sup>2</sup>, Minping Wan<sup>1</sup>, Yunke Yang<sup>1</sup>,  
Randal Burns<sup>2</sup>, Shiyi Chen<sup>1</sup>, Gregory Eyink<sup>3</sup>, Alex Szalay<sup>4</sup>

(1) Mechanical Engineering, (2) Computer Science (3) Applied Mathematics & Statistics, (4) Physics and Astronomy.

(\* Now at Sheffield Univ. UK)

Funding: National Science Foundation  
ITR and MRI




# Report on Cyber Fluid Dynamics

## Cyber-Fluid Dynamics

Report (Draft) on  
NSF Workshop on Cyber-Fluid Dynamics:  
New Frontiers in Research and Education

NSF Headquarters, Arlington VA  
July 19-20, 2007  
<http://www.nsf-cyberfluids.gatech.edu>



### Organizers:

P.K. Yeung, Georgia Institute of Technology  
R.D. Moser, University of Texas at Austin

### NSF Program in Fluid Dynamics


(in cooperation with Programs in)  
Combustion, Fire and Plasma Systems  
Dynamical Systems  
Particulate and Multiphase Processes  
Interfacial Transport and Thermodynamics  
Thermal Transport Processes

### Report Co-authors:

P.K. Yeung	Georgia Institute of Technology
R.D. Moser	University of Texas at Austin
M.W. Plesniak	Polytechnic University
C. Meneveau	The Johns Hopkins University
S. Elghobashi	University of California, Irvine
C.K. Aidun	Georgia Institute of Technology

# Report on Cyber Fluid Dynamics

Among recommendations:



.....

a culture of open communication and community-wide standards so that as many researchers as possible, including students exposed to the national supercomputing landscape, will benefit without duplication of effort. We call for leading data authors and data users in several areas to formulate community agreements on data formats and download or transfer protocols, especially for very large datasets of either computational or experimental origin. Efforts at building virtual organizations incorporating these elements, and more, are highly encouraged. We also recommend that the community work more closely, within itself and with NSF program

.....

## Motivation:


- DNS of turbulence generating huge multi-Tbyte datasets
- Analyzed on the fly during simulation by practitioner
- Must redo runs for each new question that was not foreseen during simulation,
- Or: store time-history
- If time-history is stored, one must retrieve data from large storage to large CPU for analysis, e.g. with local accounts on HPC machine.

Must know how to use and run programs there, privileges for access to large chunks of memory, casual initial analysis not possible without investing significant effort to set up such access...

## Motivation:

- This prevailing approach has hampered wide accessibility of large CFD datasets.

e.g. mathematician wishing to test whether material loops in turbulence develop fractal wrinkling will have a hard time exploring this question on (e.g.) a  $1,024^3$  DNS simulation without asking the simulator to do it for him/her.



- The fact is:

*“very large simulations remain out of reach of most”*

The problem will not automatically get better - even if “ftp” gets faster, size of “top-ranked simulations” growing even faster: **i.e. without changing current approach, top-ranked simulations will be accessible only to a shrinking subset of the scientific community.**

## Motivation:

- Develop **database techniques** for analysis of DNS
- What does that mean?

**Look to other fields where database technology  
is more developed (e.g. astronomy)**



data: Sloan Digital Sky Survey  
images: Blanton, Finkbeiner, Hogg, Schlegel, Wherry

## **Example of web-accessible astronomy database (co-developed at JHU):**

### **Sloan Digital Sky Survey (SDSS):**

- When completed: images covering more than a quarter of the sky, and a 3-D map of about a million galaxies and quasars.
- SDSS (“the Cosmic Genome Project”) is the most cited astronomical resource in the world,
- Led to more than 1,300 refereed publications so far.
- Its 2.5 Terabyte archive has been accessed more than 200 million times over the last five years.
- Some breakthroughs this astronomy resource has enabled:
  - first precision measurements of the fundamental cosmological parameters
  - discovery of the most distant quasars.



## Example of web-accessible astronomy database (co-developed at JHU):

SkyServer Schema Browser

http://cas.sdss.org/astrodr6/en/help/browser/browser.asp

News ▾ New ME webpage AstroToday POSEIDON TurbulenceResearch... ISIS GoogleScholar ISI-science JFM-AssocEditor T&F-JoT-ChiefEditor CATS Informaworld-JoT T&F

Sloan Digital Sky Survey / SkyServer

Home Tools SQL Search Schema Finding Chart Download Projects DR6 DAS Site Search Help

### Schema Browser

Glossary  
Algorithms

Search for

- + Tables
- + Views
- + **Functions**
- + Procedures
- + Constants
- + Indices

<b>fGetNearestObjIdAllEq</b>	Returns the objId of nearest photo object within @r arcmins of ra, dec
<b>fGetNearestObjIdEq</b>	Returns the objId of nearest photoPrimary within @r arcmins of ra, dec
<b>fGetNearestObjIdEqMode</b>	Returns the objId of nearest @mode PhotoObjAll within @r arcmins of ra, dec
<b>fGetNearestObjIdEqType</b>	Returns the objID of the nearest photPrimary of type @t within @r arcmin
<b>fGetNearestObjXYZ</b>	Returns nearest primary object within @r arcminutes of an xyz point (@nx,@ny, @nz).
<b>fGetObjectsEq</b>	A helper function for SDSS cutout that returns all objects within a certain radius of an (ra,dec)
<b>fGetObjectsMaskEq</b>	A helper function for SDSS cutout that returns all objects within a certain radius of an (ra,dec)
<b>fGetObjFromRect</b>	Returns table of objects inside a rectangle defined by two ra,dec pairs. Note the order of the parameters: @ra1, @ra2, @dec1, @dec2
<b>fGetObjFromRectEq</b>	Returns table of objects inside a rectangle defined by two ra,dec pairs. Note the order of the parameters: @ra1, @dec1, @ra2, @dec2
<b>fGetPhotoApMag</b>	Returns the r-band aperture magnitude and error for the given aperture.
<b>fGetUrlExpEq</b>	Returns the URL for an ra,dec, measured in degrees.
<b>fGetUrlExpId</b>	Returns the URL for an Photo objID.
<b>fGetUrlFitsAtlas</b>	Get the URL to the FITS file of all atlas images in a field
<b>fGetUrlFitsBin</b>	Get the URL to the FITS file of a binned frame given FieldID and band

SkyServer DR6 Search Form

http://cas.sdss.org/dr6/en/tools/search/form/form.asp

Mouseover the ? links for quick help, or click for the User Guide

Getting Started  
Famous places  
Get Images  
Scrolling sky  
Visual Tools  
Search  
- Radial  
- Rectangular  
- Search Form  
- Query Builder  
- SQL  
Object Crossid  
CasJobs

Show me **galaxies** in the region anywhere

with mag **0** and colors **0**

Remember that the magnitude scale is backward!  
For objects brighter than 18 in g, use g < 18.

for only objects with spectra

with redshifts between and

Please return:

Num. of Objects: Count Only 10 100 1000 All

Image Data: object IDs RA and Dec magnitudes

Spectral Data: redshift spectrum ID plate/mjd/fiber

SQL Query: Generate Query Edit Query Syntax

Your SQL command was:

```
select top 10 p.objid, p.ra, p.dec
from galaxy p, specobj s
where p.objid=s.bestobjid and (p.u BETWEEN 0 AND 20)
and (p.g BETWEEN 0 AND 20)
```

objid	ra	dec
587722952230174996	236.24709088	-0.4752703
587722952230175035	236.28686544	-0.51800857
587722952230175138	236.34217696	-0.46702629
587722952230175145	236.35011657	-0.59824048
587722952230175173	236.36935401	-0.57445286
587722952230240617	236.39731852	-0.49346093
587722952230240688	236.44526033	-0.50629014
587722952230306100	236.57612838	-0.55834271
587722952230502748	237.03379577	-0.44062416

SkyServer Object Explorer

http://cas.sdss.org/dr6/en/tools/explore/obj.asp

SkyServer DR6 Tools for Visual... SkyServer Object Explorer

**GALAXY** ra=173.74781796, dec=0.41921316, ObjId = 588848900446814264

Column names link to glossary entries. Move mouse over a column name to get its units.

mode	PRIMARY
status	TARGET PRIMARY OK_STRIPE OK_SCANLINE PSEGMENT RESOLVED OK_RUN GOOD SET
flags	STATIONARY MOVED BINNED1 CHILD
PrimTarget	TARGET_GALAXY
SecTarget	

u	g	r	i	z		
19.55	18.04	17.55	17.33	17.19		
err_u	err_g	err_r	err_i	err_z		
0.03	0.01	0.01	0.01	0.02		
run	rerun	camcol	field	obj	rowc	colc
756	44	4	387	56	549.4	1974.6
fiberMag_r	petroMag_r	devMag_r	expMag_r	psfMag_r	modelMag_r	
18.04	17.54	17.54	17.55	17.99	17.55	
extinction_r	petroRad_r	parentId	nChild			
0.06	1.796	588848900446814263	0			

**SpecObjID = 79597814924967936**

plate	mjd	fiberId	z	zErr	zConf	specClass	ra	dec	fiberMag_r	objId
282	51658	494	0.000	0.00006	1	STAR	173.74778	0.41924	17.81	588848900446814264

zStatus XCORR\_EMLINE  
zWarning NOT\_GAL  
PrimTarget TARGET\_GALAXY  
SecTarget  
eClass -0.031264  
emZ 0.000  
emConf  
xcZ 0.000  
xcConf 1

NED search  
SIMBAD search  
ADS search

Notes  
Save in Notes  
Show Notes


Print

## **Can same technology be used to store and access a full DNS database?**

- **Turbulence has concentrated coherent structures (like galaxies ?)**
- **Can data be stored in a way that facilitates particular queries?**
- **Can/should these be classified in some way?**
- **e.g. SQL search for “most intense vortices in region around high mean dissipation” ??**
- **Too cumbersome. We’ll show that slightly different approach is needed**
- **First, describe the dataset:**


## Dataset: *a 1,024<sup>4</sup> space-time history of isotropic turbulence*

- 1,024<sup>3</sup> DNS of forced isotropic turbulence, in  $(2\pi)^3$  periodic box
- pseudo-spectral, de-aliased (phase-shifting - Rogallo)
- rms velocity:  $u' = 0.679$
- Taylor microscale Reynolds number  $R_\lambda = 425$
- Kolmogorov time scale:  $\tau_K = 0.044$
- Kolmogorov length scale:  $\eta = 2.86\text{E-}03$
- Large eddy turnover time:  $T_L \approx 2.2$
- Simulation time-step:  $\delta t = 0.0002$
- Every 10 steps stored at  $\Delta t = 0.002$
- 1,024 times to be stored, in  $T_L$   
( $\sim 1$  integral time-scale)






# Database cluster architecture:



- **Each cluster node:**
  - Runs Windows 2003 Server and SQL Server 2005 (64-bit)
  - Data tables striped across 12-15 sets of mirrored disks (4 - 11 TB/node)
- **4 or 8 GB memory per node**
- **Code written in C# runs as a user defined procedure inside the SQL Server process using the CLR environment**

# Data generation and ingest:



**Will be  
27 Tbytes**


- The DNS is run on a Linux cluster at JHU
- Output is stored as standard C/Fortran arrays
- Files are copied onto the data cluster for ingest
- Data are reshaped into smaller cubes and inserted into the database using “BULK INSERT”
  - We read  $72 \times 1024 \times 1024$  slices into memory at a time

# Unified addressing and partitioning

- **Oct-trees and Z-curves (Morton)** are used for:

- Indexing space
- Indexing workload
- Data distribution
- Data partitioning




<b>Location</b>		
<b>X = 118</b> 01110110	<b>Y = 15</b> 00001111	<b>Z = 1</b> 0000001
<b>Morton-Order Key (binary)</b>		
000 <small>x<sub>7</sub> y<sub>7</sub> z<sub>7</sub></small>	100 <small>x<sub>6</sub> y<sub>6</sub> z<sub>6</sub></small>	100 <small>x<sub>5</sub> y<sub>5</sub> z<sub>5</sub></small>
100 <small>x<sub>4</sub> y<sub>4</sub> z<sub>4</sub></small>	010 <small>x<sub>3</sub> y<sub>3</sub> z<sub>3</sub></small>	110 <small>x<sub>2</sub> y<sub>2</sub> z<sub>2</sub></small>
110 <small>x<sub>1</sub> y<sub>1</sub> z<sub>1</sub></small>	011 <small>x<sub>0</sub> y<sub>0</sub> z<sub>0</sub></small>	
<b>Server</b>	<b>B+-Tree</b>	<b>Offset in Physical Block</b>



# Kernel of Computation & Edge Replication

- Most routines need “kernel of computation”
  - Spatial Interpolation

Lagrange polynomials 6th and 8th order





$$u(x) = \sum_{i=1}^6 l_{n-3+i}(x) u(x_{n-3+i})$$

$$l_j(x) = \frac{\prod_{i=n-2, i \neq j}^{n+3} (x - x_i)}{\prod_{i=n-2, i \neq j}^{n+3} (x_j - x_i)}$$





# Kernel of Computation & Edge Replication



- **Most routines need “kernel of computation”**
  - **Spatial Interpolation**  
Lagrange polynomials 6th and 8th order
  - **Velocity gradients**  
Based on same Lagrange polynomials
  - **Spatial averages**  
Not yet implemented
- **We replicate edge of each region**
  - Data needs met with a single read
  - 48% overhead ( $64^3$  to  $72^3$ )
  - Data “molecule” is  $72^3$  subcube
- **Time interpolation:**
  - Piecewise Cubic Hermite Interpolation:  
needs 4 t-neighbors

# Kernel of Computation & Edge Replication



- **Most routines need “kernel of computation”**

- **Spatial Interpolation**

Lagrange polynomials 6th and 8th order


- **Velocity gradients**

Based on same Lagrange polynomials

- **Spatial averages**

Not yet implemented

- **Time interpolation: Piecewise Cubic Hermite**  
Interpolation: needs 4 t-neighbors



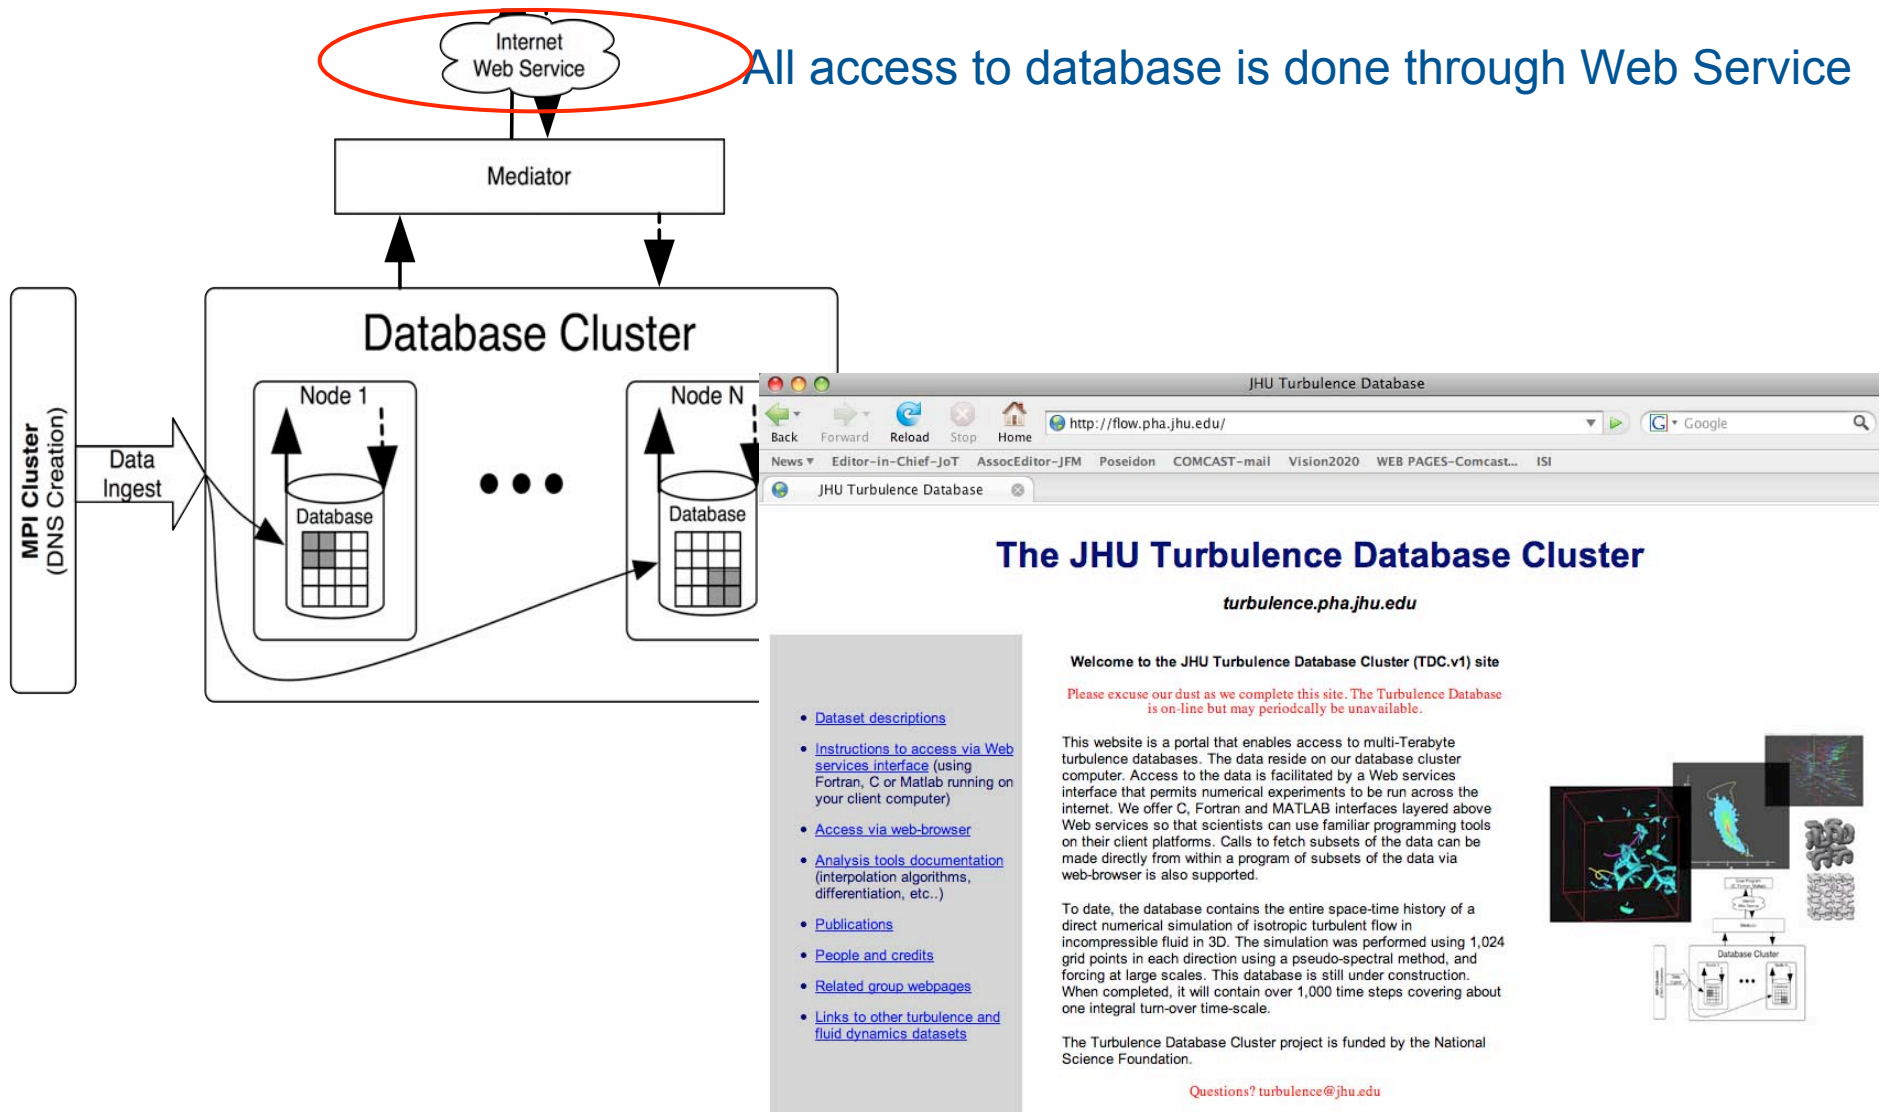
$$u(t) = a + b(t - t_n) + c(t - t_n)^2 + d(t - t_n)^2(t - t_{n+1})$$

$$a = u(t_n) \quad c = \frac{1}{\Delta t} \left( \frac{u(t_{n+1}) - u(t_n)}{\Delta t} - \dot{u}(t_n) \right)$$


$$b = \dot{u}(t_n) = \frac{u(t_{n+1}) - u(t_{n-1})}{2\Delta t}$$

$$d = \frac{2}{\Delta t^2} \left\{ \frac{1}{2} [\dot{u}(t_n) + \dot{u}(t_{n+1})] - \frac{u(t_{n+1}) - u(t_n)}{\Delta t} \right\}$$

# Accessing the database: <http://turbulence.pha.jhu.edu>



# Demo: web accessible queries:



## Query the turbulence database

This form illustrates the various operations available on the database. We will soon be updating this int

Dataset: isotropic1024coarse  
Method: GetVelocityAndPressure  
Time: 0.203 0.0 - .360  
dt = .002  
Spatial Interpolation: Lag6  
Temporal Interpolation: PCHIP  
Input method:  
☒ Single Point  
☐ Bulk - CSV File  
☐ Bulk - Binary File

Enter time

Enter position

X	Y	Z
1.132	2.3	1.3334

Go Reset

## Output

Get velocity and pressure output

X=0.8267971,Y=-0.6849948,Z=-0.0563298,P=0.3488582

## Demo: web accessible queries:

Turbulence Database Web Query

Back Forward Reload Stop Home

News Editor-in-Chief-JoT As COMCAST-mail Vision2020 WEB PAGES-Con

Turbulence Database W

Query the turbulence database

This feature allows you to query the turbulence database. We will soon be updating this interface to allow for more queries.

Query type:

Time range:

Interpolation:

Temporal Interpolation:

Input method:

☒ Single Point

☐ Bulk - CSV File

☐ Bulk - Binary File

Enter position

X	Y	Z
<input type="text" value="1.132"/>	<input type="text" value="2.3"/>	<input type="text" value="1.3334"/>

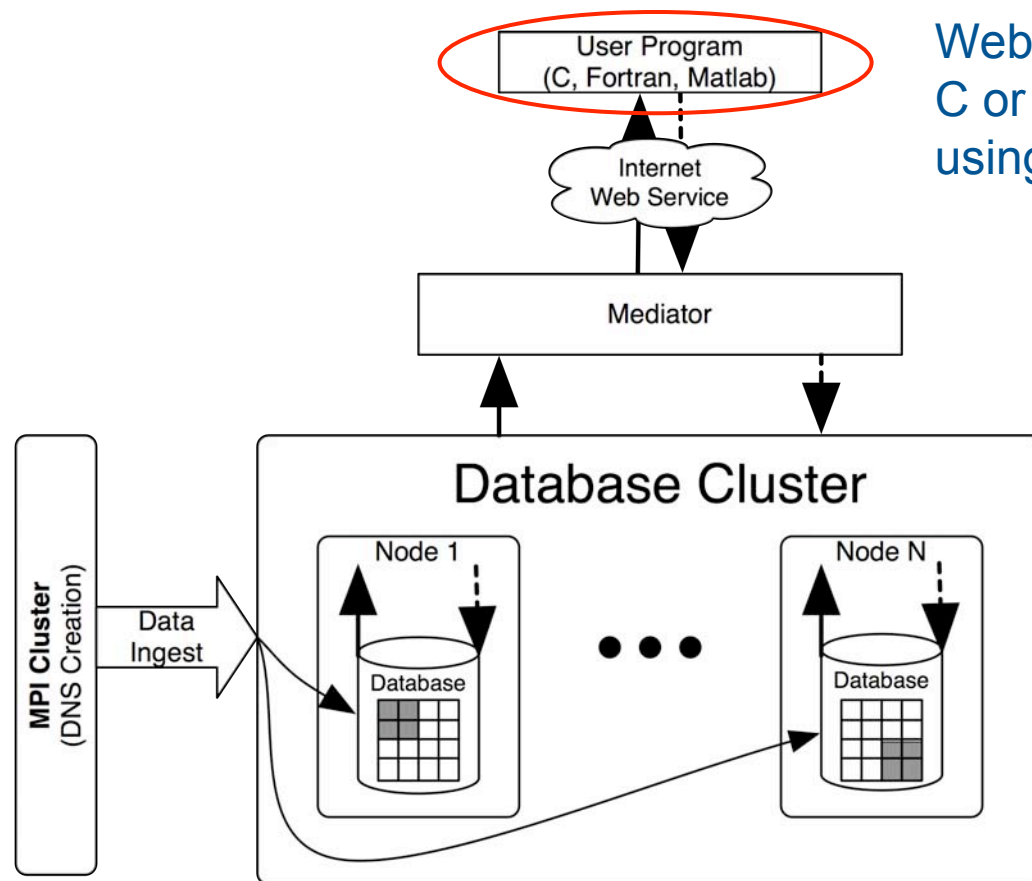
**Output**

X=0.8267971,Y=-0.6849948,Z=-0.0563298,P=0.3488582

OK for casual checks, but not useful for serious turbulence/fluids research

Get velocity and pressure output

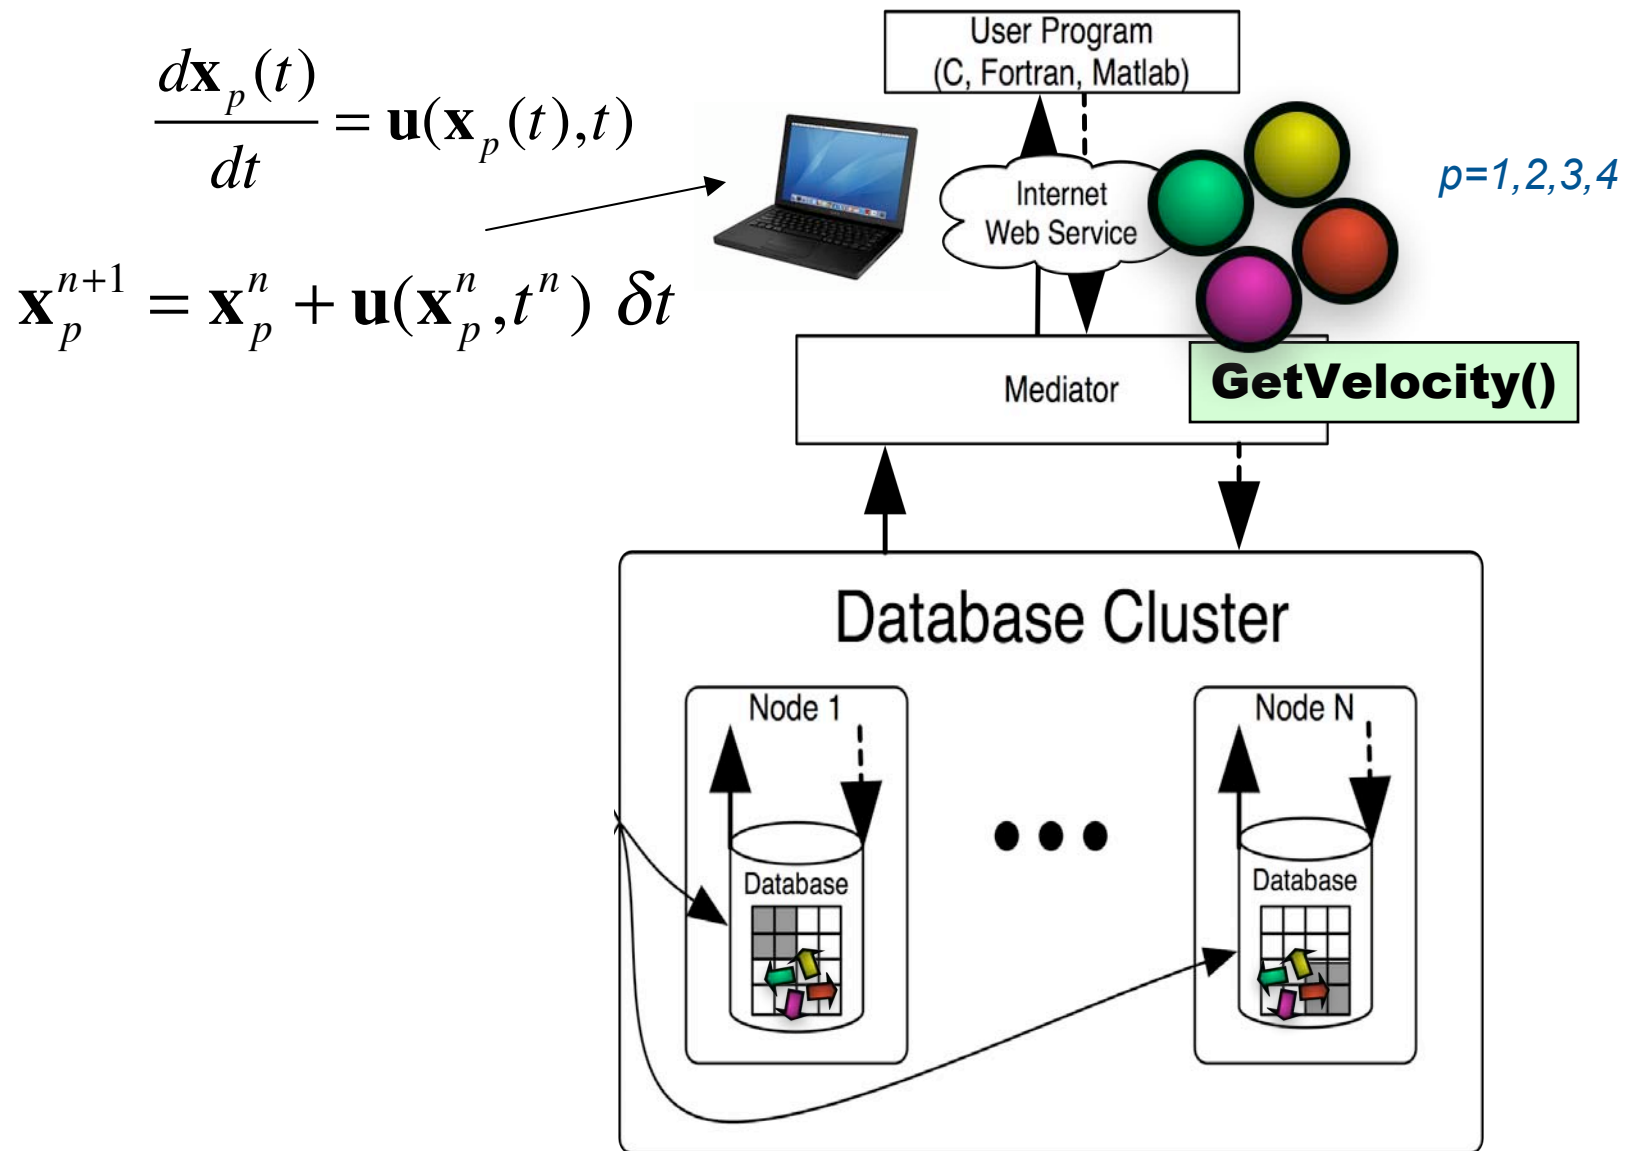
# New paradigm: client computer runs the analysis using fortran, C, matlab codes and they fetch data from database through a web-service




Web-Service calls can be made from C or Fortran on client computer using the gSOAP libraries...



# Web service request for GetVelocity



# Mediator divides workload spatially

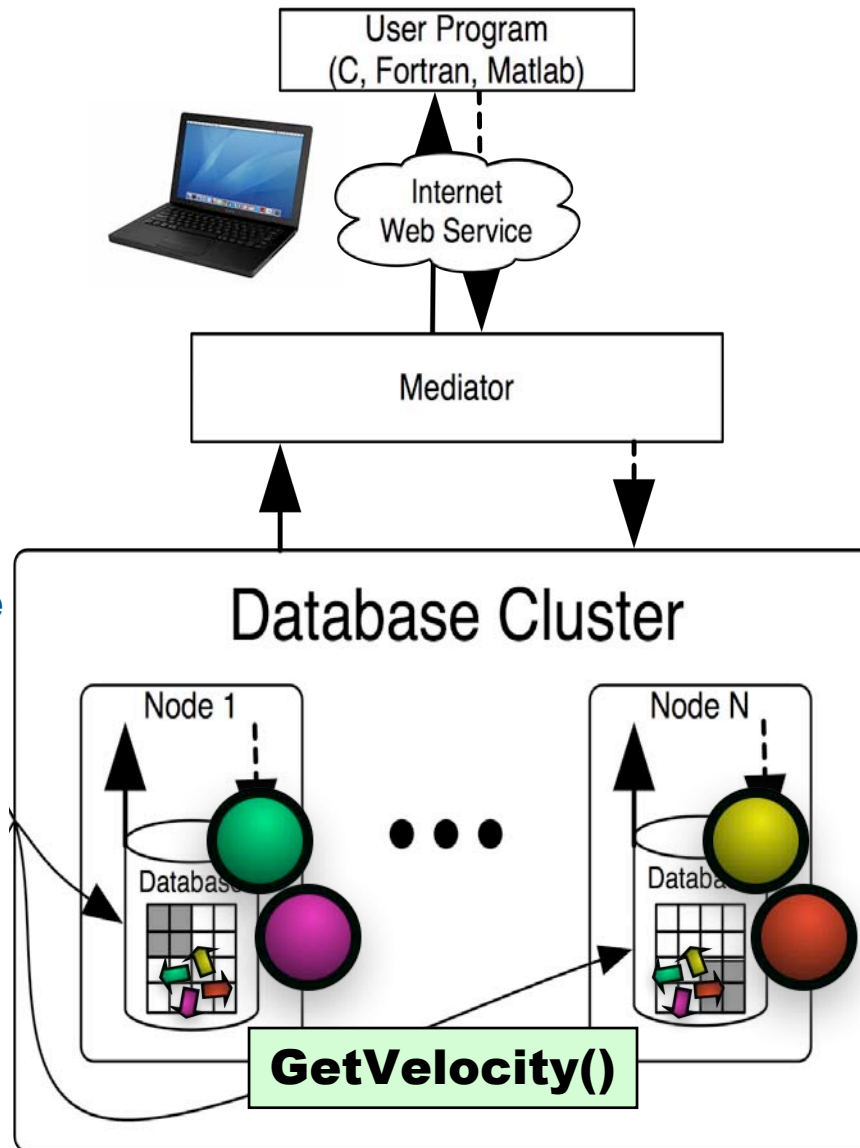




# Request sent to each database node

**Heavy local computations  
done close to the data:**

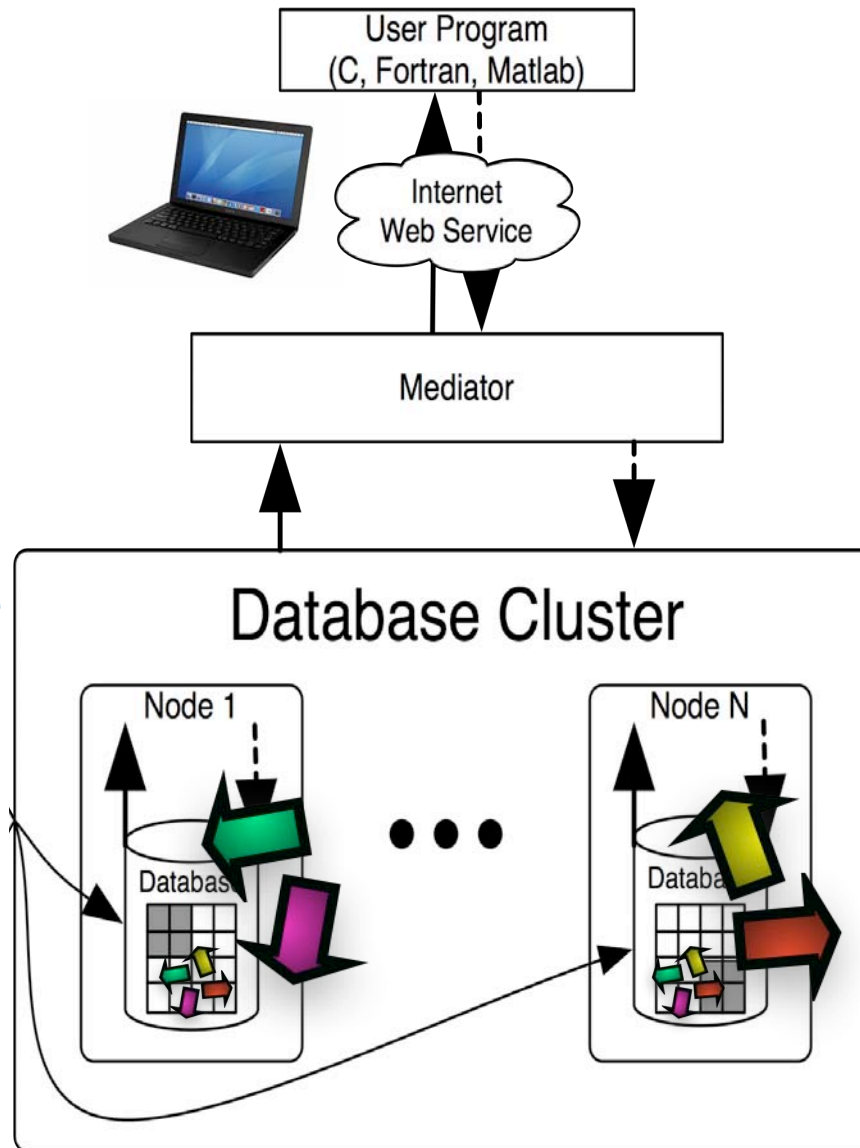
- Lagrange-polynomial interpolation in space
- P-cubic Hermite interp in time
- **Other functions:**  
vel gradient tensor, pressure, Hessian, Laplacian, pressure gradient...



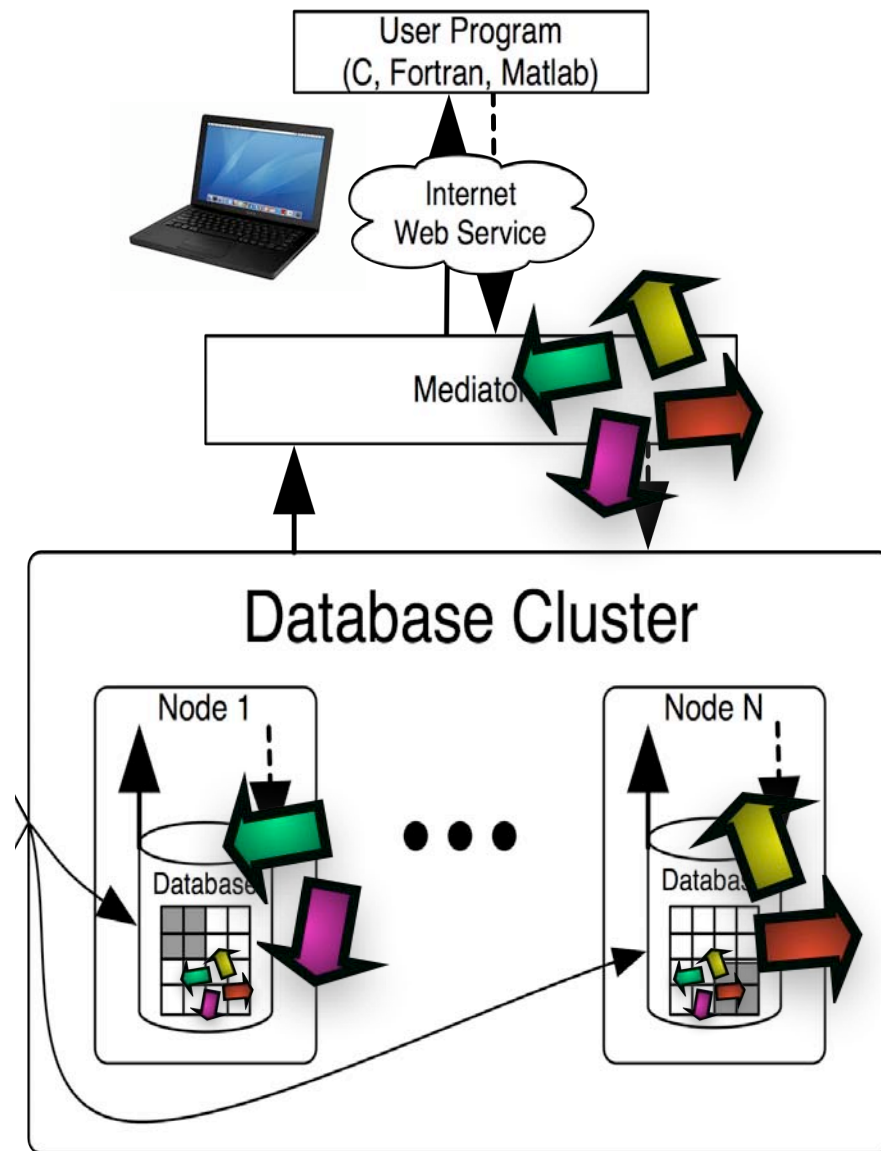
# Velocity at each location is returned

**Heavy local computations  
done close to the data:**

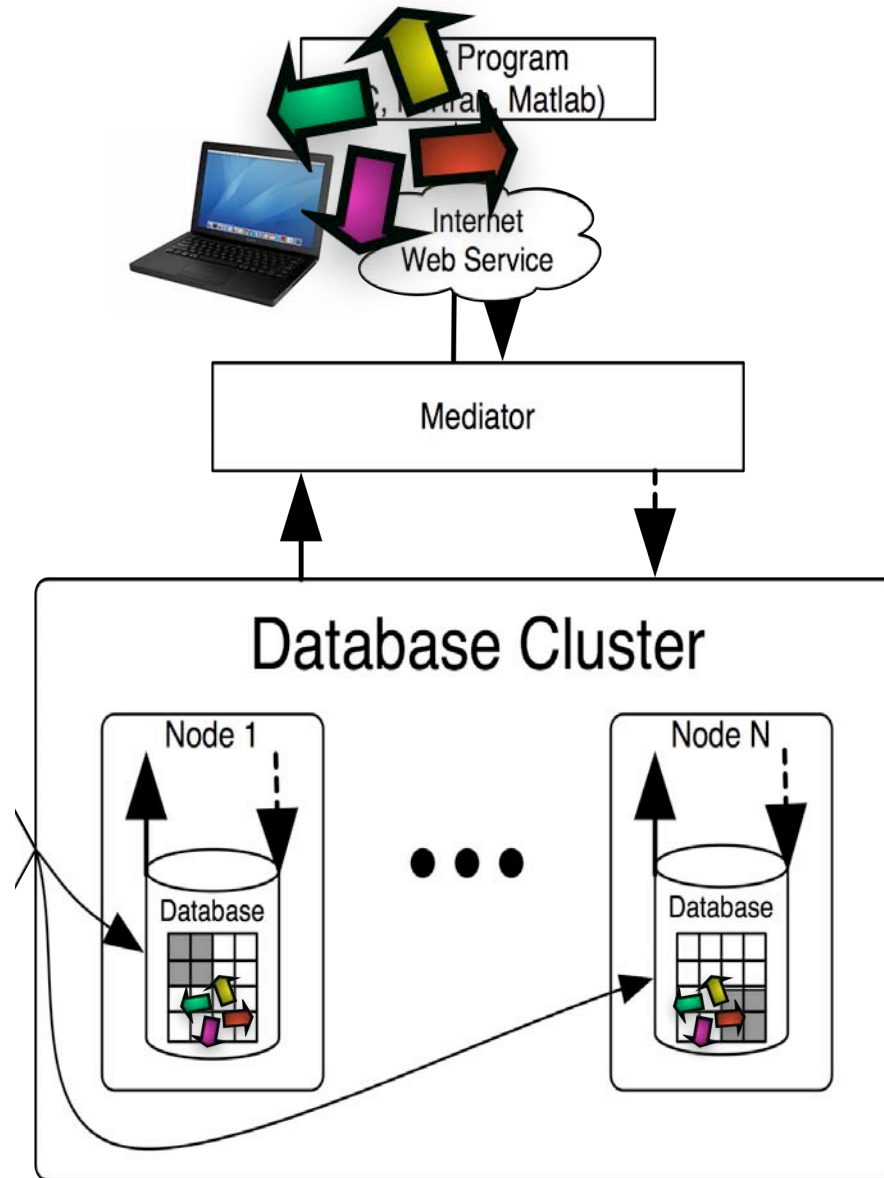
- Lagrange-polynomial interpolation in space
- P-cubic Hermite interp in time
- **Other functions:**  
vel gradient tensor, pressure, Hessian, Laplacian, pressure gradient...



# Velocity values are collated by the mediator




... and returned to the user



## Client uses velocities to determine new positions

$$\mathbf{x}_p^n + \mathbf{u}(\mathbf{x}_p^n, t^n) \delta t = \mathbf{x}_p^{n+1}$$




... and repeat with a new time  $t$ ...

# Sample code (gfortran 90) running on this Mac (unix)

```
!
!-----
! if (icase.eq.1) then
!
!   open (10,file='loop.dat',status='unknown')
!
!   write(*,*) 'A test of the JHU Turbulence Database'
!   write(*,*) 'Tracking a square loop'
!
!   time=0.15
!   deltat=0.003
!   n=250
!
!   do i=1,n,1
!     points(1,i)=1.5
!     points(2,i)=1.5
!     points(3,i)=0.5+i*0.4/n
!   end do
!   do i=1,n,1
!     points(1,i+n)=1.5
!     points(2,i+n)=1.5+i*0.4/n
!     points(3,i+n)=0.5+n*0.4/n
!   end do
!   do i=1,n,1
!     points(1,i+2*n)=1.5
!     points(2,i+2*n)=1.5+n*0.4/n
!     points(3,i+2*n)=0.5+(n-i)*0.4/n
!   end do
!   do i=1,n,1
!     points(1,i+3*n)=1.5
!     points(2,i+3*n)=1.5+(n-i)*0.4/n
!     points(3,i+3*n)=0.5
!   end do
!
!   do i=1,4*n
!     write(10,*) points(1,i),points(2,i),points(3,i)
!   end do
!   write(10,*) points(1,1),points(2,1),points(3,1)
!
!
!-----
!
```

Pick 1,000 particle positions  
around a square loop



## Sample code (gfortran 90) running on this Mac (unix)

```
!
do it = 1,15,1
!
  print *, 'time = ', time
  time=time+deltat
!
  CALL getvelocity(authkey, dataset1, time, Lagrangian6thOrder , PCHIPInterpolation, 4*n, points, dataout)
!
  do i=1,4*n
  do k=1,3
    points(k,i)=points(k,i)+dataout(k,i)*deltat
  end do
end do
!
if (it.eq.5.or.it.eq.10.or.it.eq.15) then
  do i=1,4*n
    write(10,*) points(1,i),points(2,i),points(3,i)
  end do
  write(10,*) points(1,1),points(2,1),points(3,1)
endif
write(10,*) ' '
end do
!
endif
```

advect and print-out results at subsequent times


4.3 sec. wall-time/time-step  
for 1,000 particles

# Sample code (gfortran 90) running on this Mac (unix)

```
! do it = 1,15,1
!
! print *, 'time = ', time
! time=time+deltat
!
! CALL getvelocity(authkey, dataset1, time, Lagrangian6thOrder , PCHIPInterpolation, 4*n, points, dataout)
!
! do i=1,4*n
! do k=1,3
!   points(k,i)=points(k,i)+dataout(k,i)*deltat
! end do
! end do
!
! if (it.eq.5.or.it.eq.10.or.it.eq.15) then
!   do i=1,4*n
!     write(10,*) points(1,i),points(2,i),points(3,i)
!   end do
!   write(10,*) points(1,1),points(2,1),points(3,1)
! end if
! write(10,*) ' '
! end do
!
! end if
```

advect and print-out results at subsequent times

4.3 sec. wall-time/time-  
for 1,000 particles






# Sample code (gfortran 90) running on this Mac (unix)

```
! do it = 1,15,1
!
! print *, 'time = ', time
! time=time-deltat
!
! CALL getvelocity(authkey, dataset1, time, Lagrangian6thOrder , PCHIPInterpolation, 4*n, points, dataout)
!
! do i=1,4*n
! do k=1,3
!   points(k,i)=points(k,i)+dataout(k,i)*deltat
! end do
! end do
!
! if (it.eq.5.or.it.eq.10.or.it.eq.15) then
! do i=1,4*n
!   write(10,*) points(1,i),points(2,i),points(3,i)
! end do
!   write(10,*) points(1,1),points(2,1),points(3,1)
! end if
! write(10,*) ' '
! end do
!
! end if
```

minus

advect backwards in time !

Not possible during DNS



# Demo (gfortran 90) running on this Mac (unix)

advect 10 particles with and without inertia

(Euler 1st order)

$$\frac{d\mathbf{x}_f(t)}{dt} = \mathbf{u}(\mathbf{x}_f(t), t)$$

$$\mathbf{x}_f^{n+1} = \mathbf{x}_f^n + \mathbf{u}(\mathbf{x}_f^n, t^n) \delta t$$

$$\frac{d\mathbf{v}_p(t)}{dt} = \frac{1}{\tau_p} \left( \mathbf{u}(\mathbf{x}_p(t), t) - \mathbf{v}_p(t) \right)$$

$$\mathbf{x}_p^{n+1} = \mathbf{x}_p^n + \mathbf{v}_p^n \delta t$$

# Sample code (gfortran 90) running on this Mac (unix)

```
-----
if (icase.eq.4) then
  open (50,file='dissip.dat',status='unknown')
|
write(*,*) 'A test of the JHU Turbulence Database'
write(*,*) 'extract sub-plane of velocities and dissipation for contour plot'
|
time = 0.01
n=250
|
do i=1,n
  xx(i)=0.5+2.5*i/n
  yy(i)=0.5+2.5*i/n
end do
|
do i=1,n
  do j=1,n
    ii=(i-1)*n+j
    points(1,ii)=xx(i)
    points(2,ii)=yy(j)
    points(3,ii)=1.7
  end do
end do
|
CALL getvelocitygradient(authkey, dataset1, time, Lagrangian6thOrder, NoTemporalInterpolation, n*n, points, dataout9)
|
do i=1,n
  do j=1,n
    ii=(i-1)*n+j
    vf=2.*(dataout9(2,ii)+dataout9(4,ii))*2.+2.*(dataout9(7,ii)+dataout9(3,ii))*2.+ &
      2.*(dataout9(6,ii)-dataout9(8,ii))*2.+dataout9(1,ii)**2.+dataout9(5,ii)**2.+dataout9(9,ii)**2.
    write(50,*) points(1,ii),points(2,ii),vf
  end do
  write(50,*) " "
end do
|
endif
|
-----
```

Get velocity gradients on a plane  
and evaluate **dissipation**

# Sample code (gfortran 90) running on this Mac (unix)

```
UNIX Terminal — tcsh — ttys000
-----
if (icase.eq.4) then
open (50,file='dissip.dat',status='unknown')

write(*,*) 'A test of the JHU Turbulence Database'
write(*,*) 'extract sub-plane of velocities and dissipation for contour plot'

time = 0.01
n=250

do i=1,n
  xx(i)=0.5+2.5*i/n
  yy(i)=0.5+2.5*i/n
end do

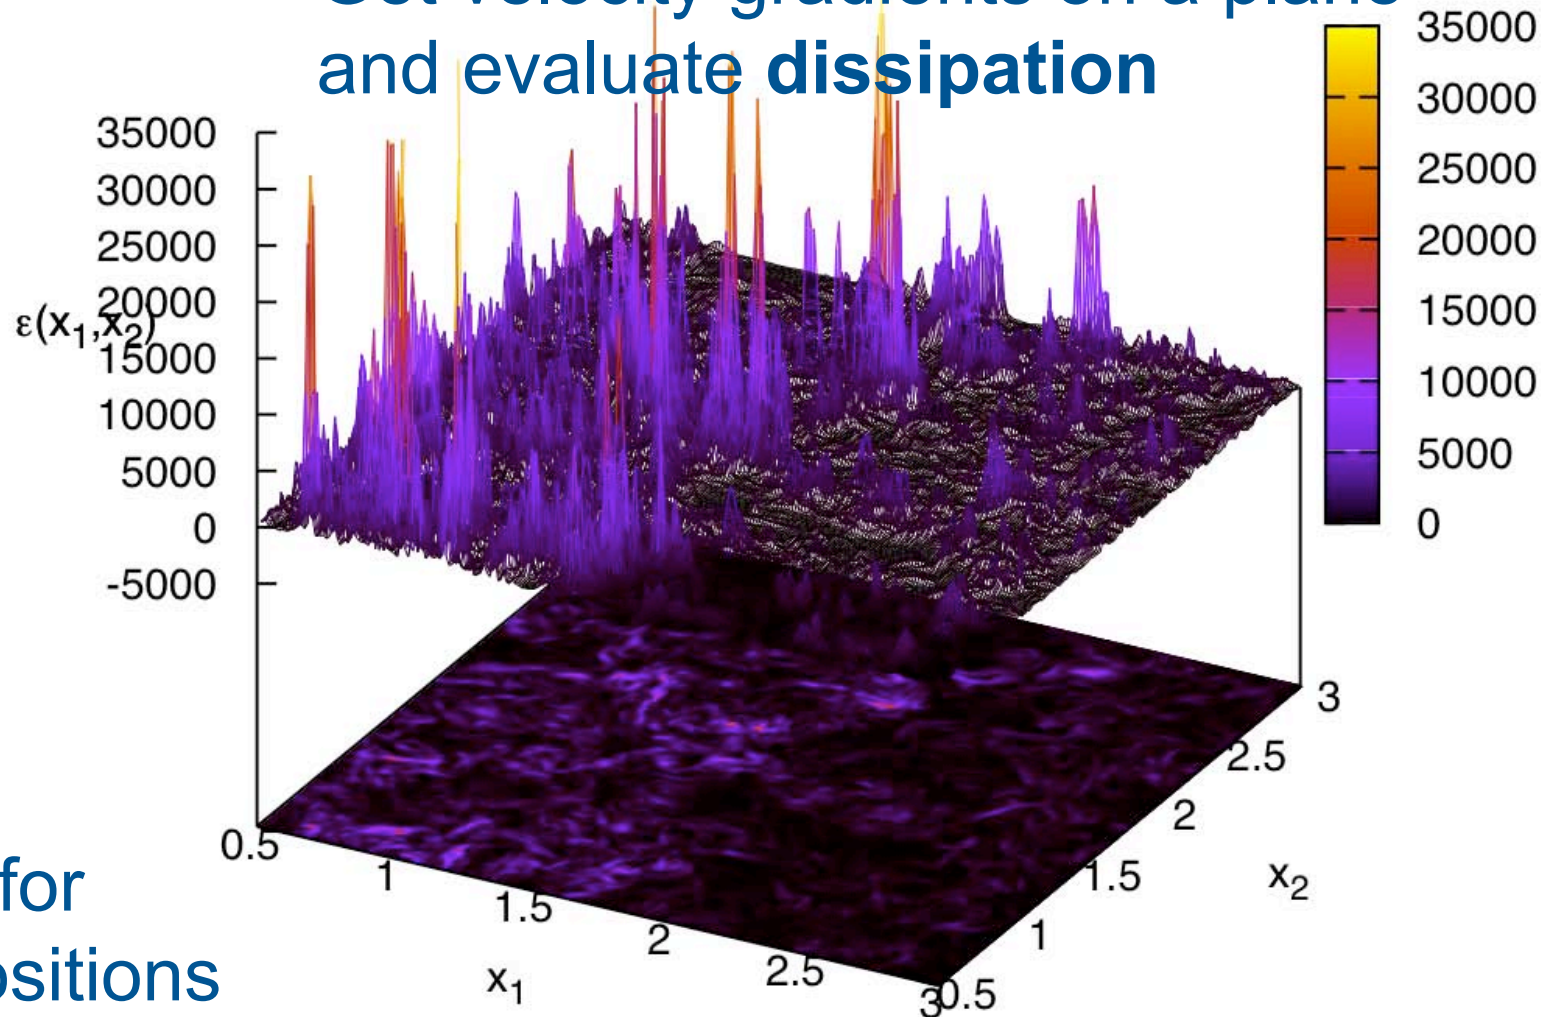
do i=1,n
  do j=1,n
    ii=(i-1)*n+j
    points(1,ii)=xx(i)
    points(2,ii)=yy(j)
    points(3,ii)=1.7
  end do
end do

CALL getvelocitygradient(

do i=1,n
  do j=1,n
    ii=(i-1)*n+j
    vf=2.*(dataout9(2,i
    2.*(dataout9(6,i
    write(50,*) points(1
  end do
  write(50,*) " "
end do

endif
-----
```

Get velocity gradients on a plane  
and evaluate dissipation



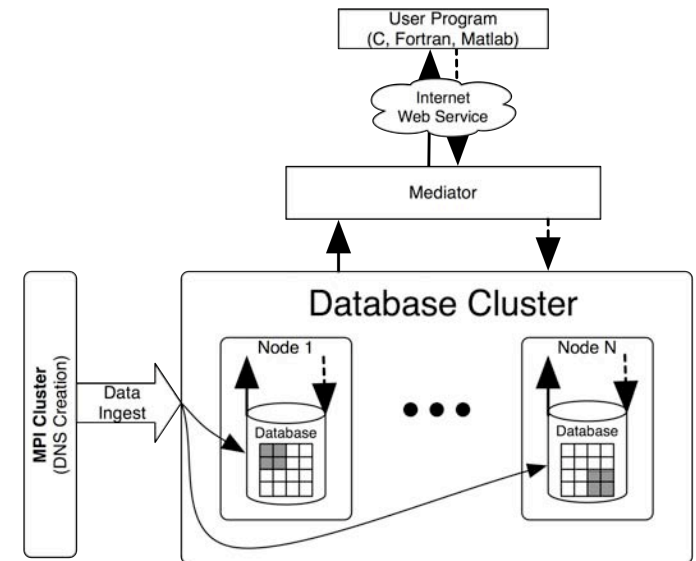
280 secs for  
62,500 positions

## Conclusions:

- Built an environment well suited for explorations of turbulent flows
- Computationally expensive DNS are performed once and then ingested and archived for possible future use by *entire* community
- Generally applicable for analysis of time-dependent, “dense”, 3-D vector and tensor fields, defined on a grid
- Brings the generic, common-place, computations (differentiation, interpolation) as close to the data as possible
- Maintains flexibility by letting user implement own computations
- But avoids having to give user accounts on server, HPC location, by providing instead a web service interfaces for Fortran, C, Matlab (using gSOAP)
- Cost-benefit: Slower analysis than when done “in core during DNS”, but much more open and easily accessible.

# Outlook:

- **Complete the time-history - 27 Terabytes**  
(as of now: “840 more time-steps to ingest”)
- **Further optimizations to accelerate service**
- **Develop additional functions, e.g. spatial filtering -**  
(but note: all-space transforms are a challenge in this approach)
- **Couple with visualization tools**
- **Include other datasets (other flows etc..)**



- **STAY TUNED AND VISIT US AT** <http://turbulence.pha.jhu.edu>